

SANKET DESHPANDE

+14042031809 | pvdsan@gmail.com | San Francisco, CA, USA | linkedin.com/in/pvdsan

EDUCATION

Master of Science, Computer Science -- Georgia State University

Aug 2023 – Aug 2025

Bachelor of Science, Computer Science -- Vellore Institute of Technology

July 2016 – June 2020

PROFESSIONAL EXPERIENCE

Digital Studio Labs

San Francisco, CA, USA

Founding ML Engineer

July 2025 - Present

- Cre8able: Built and launched a multimodal video-editing platform from idea to live beta in 6 weeks. Reduced creator editing workload by 70%, delivering clip generation via Gemini + FFmpeg on FastAPI/S3.
- Designed and deployed context-aware recommendations that combined planning context with chat history, improving adoption of suggested edits by 25%.
- Whisker: Shipped an LLM pet-care assistant in 2 weeks. Built a containerized RAG backend (Postgres + FastAPI) with sub-second response streaming; scaled to 50+ weekly active beta users.

TReNDS

Atlanta, GA, USA

Machine Learning Researcher

May 2024 - August 2025

- Built and scaled 3D ML pipelines for cognition prediction on a 12k+ subject, 300GB MRI dataset, turning raw neuroimaging into predictive insights for neuroscientists.
- Designed a custom 3D CNN with Multi-Head Self-Attention + Squeeze-and-Excitation, achieving 0.34 correlation (+15% vs. CNN baselines) and outperforming ViT and Swin Transformer approaches.
- Optimized distributed training with PyTorch DDP on a 4× GPU HPC cluster (SLURM), achieving a 7× training speedup. Ensured reproducibility with YAML configs, checkpoints, WandB, and MLFlow.
- Delivered explainable neuroimaging models using Scikit-learn, SHAP, and GradCAM, producing voxel- and region-level maps that uncovered new brain-behavior pathways.
- Published findings at IEEE ISBI (Texas, 2025) and EMBS (Denmark, 2025), validating the cerebellum's role in cognition and strengthening neuroimaging explainability research.

MORSE Studio

Atlanta, GA, USA

Computer Vision Engineer

August 2023 - April 2024

- Designed a mmWave radar perception pipeline, applying Range/Doppler FFT analysis to characterize material signatures; released open-source experiments that advanced radar-based perception.
- Developed a C-based UDP socket tool for TI mmWave radar (IWR16843 + DCA1000), enabling faster data capture and multithreaded radar stream handling — replacing proprietary tools and improving flexibility.
- Prototyped an event-driven imaging system with Luxonis neuromorphic sensors, decoding high-frequency LED signals using SciPy's Lomb-Scargle method and Python concurrency for precise real-time signal demodulation.

Harman International

Bangalore, KA, India

C++ Engineer I

October 2020 - June 2023

- Delivered production-grade C++ navigation algorithms for premium automotive OEMs (Mercedes, BMW, Audi), shipping into global vehicles and directly impacting real-time routing performance at scale.
- Engineered and deployed optimized A* navigation systems in modern C++ (C++14), resolving 350+ critical bugs and delivering 10+ major features under Agile/Scrum workflows.
- Reverse-engineered map provider APIs to classify ambiguous road link types in dense European urban networks, improving routing accuracy and navigation reliability.
- Implemented cyclic node elimination to optimize the A* search frontier in sparse North American road networks, reducing computational latency and ECU resource usage.
- Applied advanced object-oriented design, RAII, multithreading, mutex synchronization, and smart pointer-based memory management to meet strict automotive embedded performance requirements.

KPIT

Pune, MH, India

Computer Vision Intern

May 2018 - July 2018

- Built a real-time sensor-fusion pipeline (Raspberry Pi, camera, Arduino, CAN) on the car to stream multi-modal data to the cloud, enabling CV training and saving 5 hrs/week
- Developed an offline tracking tool used YOLOv2 for frame-by-frame detection, integrating it with KCF and Kalman Filter trackers for car/pedestrian labeling, thereby eliminating manual annotation and saving 2 hrs/dataset

PUBLICATIONS

Fronto-Thalamo-Cerebellar Circuitry in Predicting Cognition and Behavior of ABCD Adolescents
ISBI 2025 (1-page abstract) [\[View\]](#) EMBS 2025 (4-page paper) [\[View\]](#)

PROJECTS

Multi Agent Debate for Movie Consensus

- Implemented a multi-agent debate framework where LLM agents with distinct personas engaged in adversarial yet cooperative dialogue to refine reasoning chains and reach a consensus recommendation.
- Adopted an LLM-as-Judge evaluation loop that performed verbalized self-critique and argument scoring, yielding a significant improvement in coherence and factual alignment over single-agent baselines.
- Engineered RAG pipeline with ChromaDB vector search and Gemini 1.5 Flash for personalized recommendations, combining semantic retrieval with LLM reasoning for context-aware movie suggestions
- [Link to Project](#)

Visual Py-SLAM Toolkit

- Developed a fully end-to-end modular SLAM pipeline from raw video frames to trajectory visualization, providing a comprehensive educational framework for Computer Vision students and researchers to understand, experiment with, and extend core SLAM components without handling complex integration challenges.
- Implemented technical components including ORB feature detection/matching, essential matrix decomposition, RANSAC-based outlier rejection, and scale initialization to accurately estimate camera trajectories with robust performance on challenging indoor sequences
- [Link to Project](#)

Statistical Machine Learning for Fronto-Cerebellar Circuitry in Cognitive Prediction

- Neuroimaging Processing: Preprocessed sMRI and dMRI data from the ABCD dataset (N=10,000) using SPM12, DARTEL, and FSL. Extracted GMV and FA features, reducing preprocessing time by 20% through pipeline optimization.
- Predictive Modeling: Built Bayesian Ridge Regression, SVR, and Neural Networks for cognitive score prediction. Bayesian optimization improved SVR r^2 by 15% over default settings.
- Evaluation Metrics: Applied nested 5x5 cross-validation for robust evaluation. Achieved r^2 of 0.065 ± 0.009 for working memory using fronto-cerebellar-thalamo-parietal circuitry, outperforming fronto-parietal models.
- Comparative Insights: Conducted paired t-tests ($p < 0.05$), showing 12% higher correlation in predicting cognition with fronto-cerebellar features, highlighting the cerebellum's role in attention and working memory.
- [Link to Project](#)

Detection of milli-sized objects mmWave Radar

- Millimeter-Wave Object Detection: Detected millimeter-sized objects using 60 GHz IWR1683ISK FMCW mmWave radar hardware, leveraging MATLAB Radar Toolbox and Python libraries for advanced data processing.
- Precision Experiments: Detected penny-sized objects at a distance of 2m, surpassing the documented maximum range in controlled environments, utilizing MmWave Studio and Lua scripting for precise hardware control.
- Signal Analysis: Analyzed Range and Doppler FFTs to identify detection signatures across 6 materials and 5 distances using advanced signal processing techniques.
- [Link to Project](#)

Pose-Assisted TrackFormer: Enhanced End-to-End MOT Model

- Enhanced the TrackFormer model, developed by TUM and Facebook, by integrating pose estimation to improve identity consistency and tracking robustness.
- Developed a Pose Encoding Module using pretrained Keypoint R-CNN to extract (17 x 3) pose features (keypoints with spatial and visibility confidence), which were fused with track embeddings via a learnable linear fusion layer.
- Extended the DETR (Detection Transformer)-based architecture to incorporate pose data, optimizing detection, tracking, and association of tracks with Hungarian Matching, enabling robust re-identification even under occlusion.
- Trained on the MOT17 and CrowdHuman datasets using a two-step pipeline, evaluating performance improvements with MOTA and IDF1 metrics
- [Link to Project](#)

SKILLS

Skills: Python, OpenCV, Pytorch, Tensorflow, Keras, NumPy, Pandas, MATLAB, C/C++, Jupyter, Linux/Unix, Open3D, 3D Reconstruction, SLAM, Object Tracking, Airflow, AWS, Computer Vision, CUDA, Data Analysis, Data Structures & Algorithms, FastAPI, Kafka, Redis, Docker, Data Science, Flask, Git, LLM, Machine Learning, Neural Networks, Scikit-learn, Seaborn, Signal Processing, Radar, AWS Sagemaker